

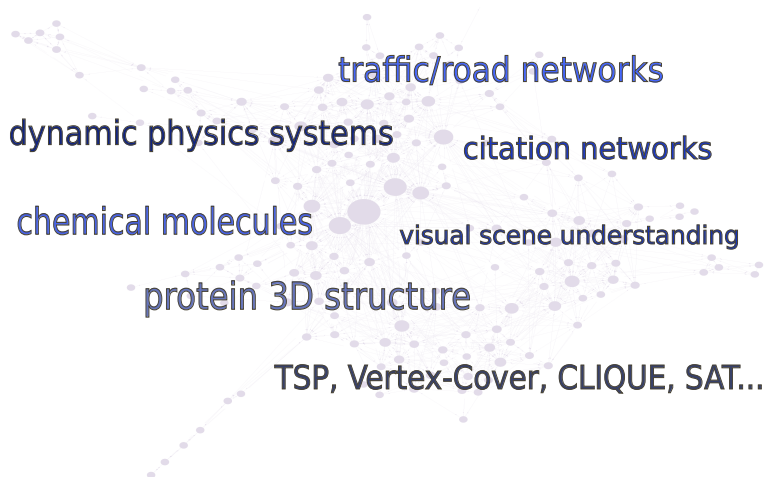
# Graph Neural Networks

Felix Becker

University of Greifswald


June 28, 2021

# Graphs are everywhere




# Inductive bias

- Often, data does not come as individual objects...



# Inductive bias


- Often, data does not come as individual objects...



- ... but rather in sets of objects and rules on how they interact.

# Inductive bias


- Often, data does not come as individual objects...



- ... but rather in sets of objects and rules on how they interact.
- Inductive bias: constraints imposed on the set of possible pairwise interactions (represented as a graph).

# Inductive bias

- Often, data does not come as individual objects...



- ... but rather in sets of objects and rules on how they interact.
- Inductive bias: constraints imposed on the set of possible pairwise interactions (represented as a graph).
- Making predictions requires 'relational reasoning' based on the graph structure.

# Inductive bias

Inductive biases can be well defined independent of the data examples:


# Inductive bias

Inductive biases can be well defined independent of the data examples:

## Sequences

>1j46\_A

ACGTAAAGTGTAAG (...)






# Inductive bias

Inductive biases can be well defined independent of the data examples:



## Sequences

>1j46\_A

ACGTAAGTGTAAG (...)



## Images



# Inductive bias

- We have dedicated models for some cases:

# Inductive bias

- We have dedicated models for some cases:
  - Recurrent architectures for sequences.

# Inductive bias

- We have dedicated models for some cases:
  - Recurrent architectures for sequences.
  - Convolutional neural networks for images.

# Inductive bias

- We have dedicated models for some cases:
  - Recurrent architectures for sequences.
  - Convolutional neural networks for images.
- But how to handle data with less well defined inductive biases?  
(e.g. chemical molecules, road networks, citation networks...)

# Graph Neural Networks (GNNs)

## Definition: Feature graph

A (directed) feature graph is a 3-tuple  $G = (u, V, E)$  with a global attribute  $u$ , nodes  $V = \{v_i\}_{i=1,\dots,n}$  where  $v_i$  are the attributes of the node at index  $i$  and edges  $E = \{(e_j, s_j, r_j)\}_{j=1,\dots,m}$  with edge attributes  $e_j$ , a sender node index  $s_j$  and a receiver node index  $r_j$ .

# Graph Neural Networks (GNNs)

## Definition: Feature graph

A (directed) feature graph is a 3-tuple  $G = (u, V, E)$  with a global attribute  $u$ , nodes  $V = \{v_i\}_{i=1,\dots,n}$  where  $v_i$  are the attributes of the node at index  $i$  and edges  $E = \{(e_j, s_j, r_j)\}_{j=1,\dots,m}$  with edge attributes  $e_j$ , a sender node index  $s_j$  and a receiver node index  $r_j$ .


## Definition: Graph neural network

A graph neural network (GNN) is a mapping  $\omega : G \mapsto G'$  that maps a feature graph  $G = (u, V, E)$  to another feature graph  $G' = (u', V', E')$  with  $V' = \{v'_i\}_{i=1,\dots,n}$  and  $E' = \{(e'_j, s_j, r_j)\}_{j=1,\dots,m}$ .

# An implementation of $\omega$

Let  $\phi_v, \phi_e, \phi_u$  be learnable non-linear functions (e.g. multilayer perceptrons).

- 1 Update all edges






# An implementation of $\omega$

Let  $\phi_v, \phi_e, \phi_u$  be learnable non-linear functions (e.g. multilayer perceptrons).


- 1 Update all edges
- 2 Aggregate neighborhoods



# An implementation of $\omega$

Let  $\phi_v, \phi_e, \phi_u$  be learnable non-linear functions (e.g. multilayer perceptrons).


- 1 Update all edges
- 2 Aggregate neighborhoods
- 3 Update all nodes



# An implementation of $\omega$

Let  $\phi_v, \phi_e, \phi_u$  be learnable non-linear functions (e.g. multilayer perceptrons).


- 1 Update all edges
- 2 Aggregate neighborhoods
- 3 Update all nodes
- 4 Aggregate nodes



# An implementation of $\omega$


Let  $\phi_v, \phi_e, \phi_u$  be learnable non-linear functions (e.g. multilayer perceptrons).

- 1 Update all edges
- 2 Aggregate neighborhoods
- 3 Update all nodes
- 4 Aggregate nodes
- 5 Update global attribute




# Message passing

Message passing: a composition of a GNN  $\omega$  with itself for a fixed number of iterations:  $\omega(\omega(\dots\omega(G)))$



# Message passing

Message passing: a composition of a GNN  $\omega$  with itself for a fixed number of iterations:  $\omega(\omega(\dots\omega(G)))$



- If the global attribute  $u$  is excluded, the output of a node  $v$  after  $N$  iterations is conditioned on all nodes with a distance of at most  $N$  to  $v$ .

# Representation learning

What do the learned representations  $v_i$ ,  $e_j$  and  $u$  mean?

# Representation learning

What do the learned representations  $v_i$ ,  $e_j$  and  $u$  mean?

Short answer: We do not know.



# Representation learning

What do the learned representations  $v_i$ ,  $e_j$  and  $u$  mean?

Short answer: We do not know.



But we can have an idea:

# Representation learning

What do the learned representations  $v_i$ ,  $e_j$  and  $u$  mean?

Short answer: We do not know.

But we can have an idea:




We use a  $D$  dimensional latent space where each neuron could represent a rather simple, independent property. Some of these properties, we can try to interpret as a human.

⇒ Empirical: A shallow model on top (e.g. a linear combination) could make accurate predictions.

# Remarks

A GNN is invariant to graph isomorphism, if the aggregation operations are symmetric functions (sum, average...)



- A GNN is differentiable, if  $\phi_v, \phi_e, \phi_u$  are (w.r.t. their weights  $\theta$ ).

# Remarks

- A GNN is differentiable, if  $\phi_v, \phi_e, \phi_u$  are (w.r.t. their weights  $\theta$ ).
- Therefore, we can backpropagate a loss signal in order to update  $\theta$ .

# Remarks

- A GNN is differentiable, if  $\phi_v, \phi_e, \phi_u$  are (w.r.t. their weights  $\theta$ ).
- Therefore, we can backpropagate a loss signal in order to update  $\theta$ .
- The loss may depend on  $u'$  (graph focused),  $V'$  (node focused) or  $E'$  (edge focused) or all of these.

# Some GNN applications

## Predict properties of chemical molecules

Message Passing Neural Net



Neural Message Passing for Quantum Chemistry

– Gilmer et al., 2017

## Input

Molecule as a graph: Type of atom for each node, type of bond for each edge.

# Some GNN applications

## Predict properties of chemical molecules

### Message Passing Neural Net



Neural Message Passing for Quantum Chemistry

– Gilmer et al., 2017

## Input

Molecule as a graph: Type of atom for each node, type of bond for each edge.

## Output


A global property of the molecule e.g. energy



# Some GNN applications

## Predict properties of chemical molecules

### Message Passing Neural Net



### Neural Message Passing for Quantum Chemistry

– Gilmer et al., 2017

## Input

Molecule as a graph: Type of atom for each node, type of bond for each edge.

## Output

A global property of the molecule e.g. energy


## Loss

MSE on the global output  $u'$

# Some GNN applications

## Solve SAT

Learning a SAT solver from  
single-bit supervision  
– Selsam et al., 2019




## Input

SAT instance  $I$  encoded as an unlabeled graph  $G$

# Some GNN applications

## Solve SAT

Learning a SAT solver from  
single-bit supervision  
– Selsam et al., 2019



## Input

SAT instance  $I$  encoded as an unlabeled graph  $G$


## Output

$P(I \text{ is satisfiable} \mid G)$

# Some GNN applications

## Solve SAT

Learning a SAT solver from  
single-bit supervision  
– Selsam et al., 2019



## Input

SAT instance  $I$  encoded as an unlabeled graph  $G$

## Output


$P(I \text{ is satisfiable} \mid G)$

## Loss

Binary cross entropy on the global output  $u'$

# Some GNN applications

## Timesteps in a dynamic physics system



([github.com/deepmind/graph\\_nets](https://github.com/deepmind/graph_nets))

Relational inductive biases, deep learning, and graph networks  
– Battaglia et al., 2018

## Input

Initial state of the rope as a graph: Masses and positions as node attributes. Only add edges for adjacent masses on the rope.  
Possible global attribute: Gravity

# Some GNN applications

## Timesteps in a dynamic physics system



([github.com/deepmind/graph\\_nets](https://github.com/deepmind/graph_nets))

Relational inductive biases, deep learning, and graph networks  
– Battaglia et al., 2018

## Input


Initial state of the rope as a graph: Masses and positions as node attributes. Only add edges for adjacent masses on the rope.  
Possible global attribute: Gravity

## Output

State of the rope (i.e. updated node positions) after  $N$  timesteps.

# Some GNN applications

## Timesteps in a dynamic physics system



([github.com/deepmind/graph\\_nets](https://github.com/deepmind/graph_nets))

Relational inductive biases, deep learning, and graph networks  
– Battaglia et al., 2018

## Input

Initial state of the rope as a graph: Masses and positions as node attributes. Only add edges for adjacent masses on the rope.  
Possible global attribute: Gravity

## Output

State of the rope (i.e. updated node positions) after  $N$  timesteps.

## Loss

MSA on the node outputs  $v_i'$

# A practical example: Sorting numbers

## Input

A list of pairwise distinct numbers  $x_1, \dots, x_k$

- WLOG  $x_i \in [0, 1]$
- $k$  is variable

## Output

A sorted list of the same numbers



# A practical example: Sorting numbers

## Input

A list of pairwise distinct numbers  $x_1, \dots, x_k$

- WLOG  $x_i \in [0, 1]$
- $k$  is variable

## Output

A sorted list of the same numbers

Let  $\pi$  be a sorting permutation of the indices, i.e.  $\pi(i) = j \Leftrightarrow x_i$  is the  $j$ th smallest element

# A naive approach without GNNs

We could try to solve the problem by using a simple neural network with an input vector of fixed length  $L$ :

$$(x_1, x_2, \dots, x_k, P, P, \dots, P)$$

where  $L$  is large enough for all reasonable inputs and  $P$  is some *padding* added to input lists shorter than  $L$ .

# A naive approach without GNNs

We could try to solve the problem by using a simple neural network with an input vector of fixed length  $L$ :

$$(x_1, x_2, \dots, x_k, P, P, \dots, P)$$

where  $L$  is large enough for all reasonable inputs and  $P$  is some *padding* added to input lists shorter than  $L$ .


Why **NOT** do it this way.

- maximum list length  $L$  is not intuitive
- overhead for  $k \ll L$
- no generalization to longer lists than seen during training

# Idea

- 1 Construct a fully connected graph with numbers as nodes
- 2 Refine nodes/edge embeddings with message passing
- 3 Decode  $P(\pi(j) = \pi(i) + 1 | x_1, x_2, \dots, x_k)$  for all  $i, j = 1, \dots, k$


FULLY CONNECTED  
INPUT GRAPH



$2 \times N$   
MESSAGE PASSING



DECODING



# Loss and Training

For details see notebook "Sort.ipynb".